

CrossingFlow: Minimizing Domain Crossings in Heterogeneous Systems Under Error Constraints

Jessy Morissette

Working Paper — February 2026

Abstract

Heterogeneous compute systems—combining analog, digital, near-memory, multi-chiplet, and multi-voltage-domain substrates—promise order-of-magnitude efficiency gains. However, these gains are systematically eroded by **domain crossings**: the overhead incurred when data transitions between domains. We formalize a cost decomposition separating intra-domain compute from crossing cost, introduce elasticity ϵ as a dominance indicator, and verify the resulting Domain Crossing Law across five boundary classes and five technology nodes (3–28nm). Crossing-to-compute cost ratios range from 5 \times to 32,000 \times (geometric mean: 73 \times in pJ/byte). Elasticity reaches $\epsilon \approx 0.86\text{--}1.00$ in crossing-dominated regimes, confirming near-linear energy reduction from crossing elimination. We present CrossingFlow, a compiler that treats crossings as a scarce resource and minimizes crossing volume under error and latency constraints.

1. Introduction

The end of Dennard scaling has driven architecture toward heterogeneous systems combining multiple compute substrates. Analog compute-in-memory (CIM) promises sub-femtojoule MACs; chiplet architectures enable specialized die integration; near-memory processing reduces data movement. Each achieves impressive intra-domain efficiency. Yet system-level gains consistently fall short of component-level projections.

We argue this gap is explained by a single phenomenon: the cost of **domain crossings**. When data crosses any boundary—ADC/DAC, DRAM fetch, inter-die transfer, level-shifter—the system pays 5 \times to 32,000 \times more per byte than intra-domain compute. This observation, the **Domain Crossing Law**, has three corollaries: (1) improving intra-domain efficiency yields diminishing returns once crossings dominate; (2) architectural crossing reduction gives near-linear gains; (3) the correct compiler target is crossings, not operations.

Contributions: (i) formal cost decomposition in pJ/byte; (ii) taxonomy of five boundary classes; (iii) elasticity ϵ as dominance thermometer; (iv) three invariance tests plus ablation; (v) CrossingFlow compiler with multi-domain IR and crossing-minimizing optimizer.

1.3 Distinction from Prior Work

Horowitz (2014) quantified a single boundary: memory \leftrightarrow compute. That seminal result was descriptive and boundary-specific. Our contribution is structurally different in three ways. First, we generalize from one boundary to a unified framework covering all domain crossings, with a common cost metric (pJ/byte) enabling cross-boundary comparison. Second, we introduce elasticity ϵ as a quantitative dominance indicator providing actionable guidance: optimize V_b when

ε is high, optimize c_b when ε is low. Third, we demonstrate structural invariance across boundary types and technology nodes, suggesting an architectural law rather than a technology-specific observation.

2. The Domain Crossing Law

2.1 Definitions

Domain boundary: any interface where data changes representation, location, or operating constraints. Each boundary b has crossing cost c_b (pJ/byte). **Crossing volume V_b :** total bytes traversing boundary b during execution. **Intra-domain cost C_{intra} :** energy of operations within a single domain.

2.2 Decomposition

$$C(P, H) = C_{intra}(P, H) + \sum_b V_b(P, H) \cdot c_b(H)$$

This is exact when crossing costs are independent and per-byte cost is constant (burst-amortizable). Both hold within 10–15% for all five boundary types studied.

2.3 Elasticity

$$\varepsilon_b = \frac{\partial \log(C)}{\partial \log(V_b)}$$

When crossing cost dominates, $\varepsilon \rightarrow 1$: every 1% crossing reduction yields ~1% energy reduction. When compute dominates, $\varepsilon \rightarrow 0$. The elasticity is the decision criterion: optimize V_b when ε is high; optimize c_b when ε is low.

3. Experimental Setup

We study five boundary types with costs from published sources, normalized to pJ/byte at 7nm.

Boundary	c_compute	c_crossing	Ratio	Source
Analog↔Digital	0.0001	3.20	32,000×	ISSCC 2023, Walden FoM
Memory↔Compute	0.25	1.25	5×	Horowitz 2014, 7nm scaled
Chiplet↔Chiplet	0.25	5.00	20×	PCIe 1.0 spec
Near↔Far Memory	0.25	10.00	40×	Samsung HBM3
Voltage↔Voltage	0.05	0.80	16×	Multi-Vdd literature

Table 1. Crossing-to-compute cost ratios at 7nm (pJ/byte). Geometric mean: 73×.

Analog↔Digital: 0.1 fJ/MAC (ISSCC 2023 best CIM). One MAC \approx 1 byte. ADC: Walden FoM 50 fJ/conv-step, 6-bit $\rightarrow 50 \times 2^6 = 3,200$ fJ/value = 3.2 pJ/byte. **Memory↔Compute:** Horowitz (2014) scaled to 7nm: ~1.25 pJ/byte DRAM access vs 0.25 pJ/byte compute. **Chiplet:** PCIe short-

reach ~ 0.5 pJ/bit + overhead = ~ 5 pJ/byte. **HBM:** Samsung HBM3 ~ 3.9 pJ/bit + controller ≈ 10 pJ/byte. **Level shifter:** ~ 0.1 pJ/bit = 0.8 pJ/byte.

Canonical workload: 256 KB intra-domain compute, crossing volume swept 256 B to 512 KB.

4. Results

4.1 Test A: Technology Invariance

We repeat across 3nm–28nm. Compute scales with transistor density; crossing scales more slowly (physical phenomena: charge redistribution, signal propagation).

Boundary	3nm	5nm	7nm	14nm	28nm
Analog↔Digital	56,000×	42,667×	32,000×	20,800×	12,000×
Memory↔Compute	9×	7×	5×	3×	2×
Chiplet↔Chiplet	35×	27×	20×	13×	8×
Near↔Far Memory	70×	53×	40×	26×	15×
Voltage↔Voltage	28×	21×	16×	10×	6×

Table 2. Ratios across technology nodes. Ratios decrease but remain $\gg 1$ at all nodes.

4.2 Test B: Scale Invariance

The tipping point (crossing fraction where crossing energy > 50%) ranges from 0.1% (analog, extreme sensitivity) to 50% (memory, low sensitivity). All fall within realistic workload parameters.

4.3 Test C: Elasticity

Boundary	Ratio	ϵ crossing	ϵ transition	ϵ compute
Analog↔Digital	32,000×	0.996	—	—
Memory↔Compute	5×	0.763	0.31–0.47	0.01–0.05
Chiplet↔Chiplet	20×	0.838	0.31–0.47	0.03–0.10
Near↔Far Memory	40×	0.862	0.31–0.47	0.05–0.10
Voltage↔Voltage	16×	0.865	0.26–0.42	0.02–0.08

Table 3. Elasticity ϵ by regime. In crossing-dominated regimes, $\epsilon \approx 0.86$ –1.00.

The analog↔digital boundary achieves $\epsilon = 0.996$: a 50% reduction in conversions yields 49.8% energy reduction. Even memory↔compute, with modest 5× ratio, reaches $\epsilon = 0.76$ in its crossing regime.

4.4 Ablation: Hardware vs Architecture

Comparing 10× better crossing technology vs 10× fewer crossings: at 25% crossing fraction, both yield comparable gains. The critical insight is that architectural gains compose multiplicatively

across boundaries, while technological gains are bounded by physics. In fully crossing-dominated regimes, architecture strictly wins.

4.5 Real Workload Validation

We validate the law on two production workloads: ResNet-50 inference (ImageNet, batch=1, 4.31 GMAC) and GPT-2 Small attention ($d_{model}=768$, variable sequence length). ResNet-50 is mapped to a CIM accelerator (128×128 crossbar tiles, 6-bit SAR ADC), a DRAM-backed GPU, and a 2-chiplet architecture.

Boundary	Cross frac	ϵ	Regime	Prediction
Analog↔Digital (CIM)	98.9%	0.989	CROSSING	✓ Confirmed
Memory↔Compute (DRAM)	8.1%	~0.08	COMPUTE	✓ Confirmed
Chiplet↔Chiplet	0.05%	~0.00	COMPUTE	✓ Confirmed

Table 4. ResNet-50: crossing fraction and regime per boundary. The law correctly predicts the dominant cost component in each case.

For the Transformer workload, the CIM boundary remains crossing-dominated at all sequence lengths (95.8–98.8% crossing fraction). The memory boundary is compute-dominated at all scales tested (0.3–7.6%), consistent with the Transformer’s $O(n^2)$ compute scaling outpacing its linear memory traffic. This demonstrates that the law correctly identifies both crossing-dominated and compute-dominated regimes in real workloads, and that ϵ serves as a reliable design-time predictor.

4.6 Illustrative Compilation Example

To demonstrate the practical utility of crossing-aware compilation, we map a 3-layer CIM network (256×256 per layer) under two schedules:

Baseline: ADC after each layer. $N_{cross} = 3 \times 256 = 768$ conversion events. $E_{total} = 19,661$ pJ (crossings: 99.6%).

CrossingFlow: Fuse layers 1–2 in analog domain (noise budget allows), ADC only after layer 3. $N_{cross} = 256$ conversion events. $E_{total} = 6,580$ pJ. **Energy reduction:** $3.0\times$ from crossing elimination alone.

This confirms Corollary 2: reducing N_h by $3\times$ yields approximately $3\times$ energy reduction when $\epsilon \approx 1$. The compiler’s noise budget system verified that the fused chain accumulated error remained within the 5% tolerance threshold.

5. CrossingFlow Compiler

CrossingFlow operationalizes the Domain Crossing Law. Its IR annotates every tensor with its current domain (analog-voltage, charge, time, digital). Domain transitions are explicit cast

operations with costs. A noise-budget type system tracks cumulative error through the graph. The optimizer solves:

$$\min \sum_b V_b \cdot c_b \quad \text{s.t.} \quad \text{error} \leq \tau, \quad \text{latency} \leq L$$

using three strategies: fusion (eliminate inter-layer crossings), domain extension (stay in cheaper domain longer), and batched conversion (amortize cost over accumulated results).

6. Discussion

When the law breaks down. In compute-bound workloads with high arithmetic intensity and minimal data movement ($\varepsilon \approx 0$), crossing reduction is ineffective. The law correctly identifies this regime via the elasticity metric. Such workloads are increasingly rare in modern AI applications.

Post-Moore implications. The law reframes the research agenda: instead of better components (ADCs, compute dies, memory cells), the highest-leverage investments are in reducing crossing frequency through architectural co-design and novel interface technologies (time-domain encoding, charge-domain accumulation).

Cost model refinement. Our linear model $C_{\text{cross}} = V_h \cdot c_h$ captures the dominant per-byte term. In practice, crossings include fixed overhead: $C_{\text{cross}} = \alpha_h \cdot \text{events} + \beta_h \cdot \text{bytes}$. Our model captures β_h , which dominates for typical burst sizes (≥ 32 bytes). The α_h term (ADC calibration, chiplet link setup) becomes significant only for very small transfers.

Sensitivity of the 32,000 \times ratio. The analog \leftrightarrow digital ratio uses the best published CIM crossbar efficiency (0.1 fJ/MAC). Including peripheral overhead (1.3 \times multiplier) reduces this to 24,600 \times . A 10 \times pessimistic compute assumption yields 3,200 \times . Under all perturbations, the qualitative conclusion—crossing dominance for CIM—is unchanged.

Limitations. Our model assumes linear per-byte costs and constant c_h over device lifetime. In practice, c_h varies due to conductance drift (analog), electromigration (interconnect), and HBM degradation. For time-domain encoding, the O(1)-vs-resolution advantage holds for ≤ 8 bits; at higher resolutions, jitter and counter precision may restore scaling. All results are simulation-based; silicon validation remains future work.

7. Related Work

The energy of data movement was quantified by Horowitz (2014), showing DRAM access dominates compute. This drove near-memory and processing-in-memory architectures. We generalize from one boundary type to all domain crossings. In analog CIM, the ADC bottleneck is well documented (60–80% of accelerator energy/area). ADC-less architectures (rTD-CiM, HCiM) address this at the hardware level; CrossingFlow provides the software counterpart. Heterogeneous compilers (MLIR, TVM, Halide) focus on scheduling and layout; we extend this with crossing-first optimization and noise-budget types.

8. Conclusion

We have presented the Domain Crossing Law: heterogeneous system efficiency is dominated by crossing count under error constraints. Across five boundaries and five nodes, crossing costs exceed compute by 5–32,000×. Architectural crossing elimination yields near-linear gains ($\epsilon \approx 0.86\text{--}1.00$). CrossingFlow operationalizes this as a compiler that minimizes crossing volume.

This shift reframes compiler optimization from operation scheduling to domain-transition minimization. **The correct unit of optimization in heterogeneous computing is not the operation—it is the crossing.**

References

- [1] M. Horowitz, “Computing’s Energy Problem,” ISSCC, 2014.
- [2] A. Shafiee et al., “ISAAC: A CNN Accelerator with In-Situ Analog Arithmetic,” ISCA, 2016.
- [3] UCIE Consortium, “UCIE 1.0 Specification,” 2022.
- [4] Samsung, “HBM3 Product Specification,” 2023.
- [5] B. Murmann, “ADC Performance Survey 1997–2024,” Stanford.
- [6] W. Wan et al., “Compute-in-Memory Chip Based on RRAM,” Nature, 2022.
- [7] X. Peng et al., “DNN+NeuroSim V2.0,” IEEE TCAD, 2021.